



```
In [54]: #!pip install kagglehub
```

```
In [1]: import kagglehub
```

```
# Download latest version
path = kagglehub.dataset_download("camnugent/california-housing-prices")

print("Path to dataset files:", path)
```

```
Path to dataset files: C:\Users\MOHAMMED HAYATH RK\.cache\kagglehub\datasets\camnugent\california-housing-prices\versions\1
```

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from warnings import filterwarnings
filterwarnings('ignore')
```

```
In [3]: import os
```

```
In [4]: # List files
files = os.listdir(path)
print(files)

['housing.csv']
```

```
In [5]: path
```

```
Out[5]: 'C:\\\\Users\\\\MOHAMMED HAYATH RK\\\\.cache\\\\kagglehub\\\\datasets\\\\camnugent\\\\california-housing-prices\\\\versions\\\\1'
```

```
In [6]: # Use the path returned by kagglehub
base_path = path # path from kagglehub
```

```
In [7]: housing = pd.read_csv(os.path.join(base_path, "housing.csv"))
housing.head()
```

```
Out[7]: longitude latitude housing_median_age total_rooms total_bedrooms popul
      0 -122.23    37.88            41.0       880.0        129.0
      1 -122.22    37.86            21.0      7099.0       1106.0        2
      2 -122.24    37.85            52.0      1467.0        190.0
      3 -122.25    37.85            52.0      1274.0        235.0
      4 -122.25    37.85            52.0      1627.0        280.0
```

Basic EDA

```
In [8]: housing.shape
```

```
Out[8]: (20640, 10)
```

```
In [9]: housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   longitude        20640 non-null   float64
 1   latitude         20640 non-null   float64
 2   housing_median_age 20640 non-null   float64
 3   total_rooms      20640 non-null   float64
 4   total_bedrooms   20433 non-null   float64
 5   population       20640 non-null   float64
 6   households       20640 non-null   float64
 7   median_income    20640 non-null   float64
 8   median_house_value 20640 non-null   float64
 9   ocean_proximity  20640 non-null   object  
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

```
In [10]: housing.describe()
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms
count	20640.000000	20640.000000	20640.000000	20640.000000	20433.000000
mean	-119.569704	35.631861	28.639486	2635.763081	537.343750
std	2.003532	2.135952	12.585558	2181.615252	421.737500
min	-124.350000	32.540000	1.000000	2.000000	1.000000
25%	-121.800000	33.930000	18.000000	1447.750000	296.000000
50%	-118.490000	34.260000	29.000000	2127.000000	435.000000
75%	-118.010000	37.710000	37.000000	3148.000000	647.000000
max	-114.310000	41.950000	52.000000	39320.000000	6445.000000

```
In [11]: housing.isnull().sum()
```

```
Out[11]: longitude      0
          latitude       0
          housing_median_age 0
          total_rooms      0
          total_bedrooms   207
          population       0
          households       0
          median_income    0
          median_house_value 0
          ocean_proximity  0
          dtype: int64
```

Best way to handle missing values is to use
Median Imputation

Why? Because

- Median works better when the data has outliers
- Avoids the shifting the distribution
- Keeps model performance stable

```
In [12]: # Fill missing values in total_bedrooms using median
housing['total_bedrooms'] = housing['total_bedrooms'].fillna(housing['total_be
```

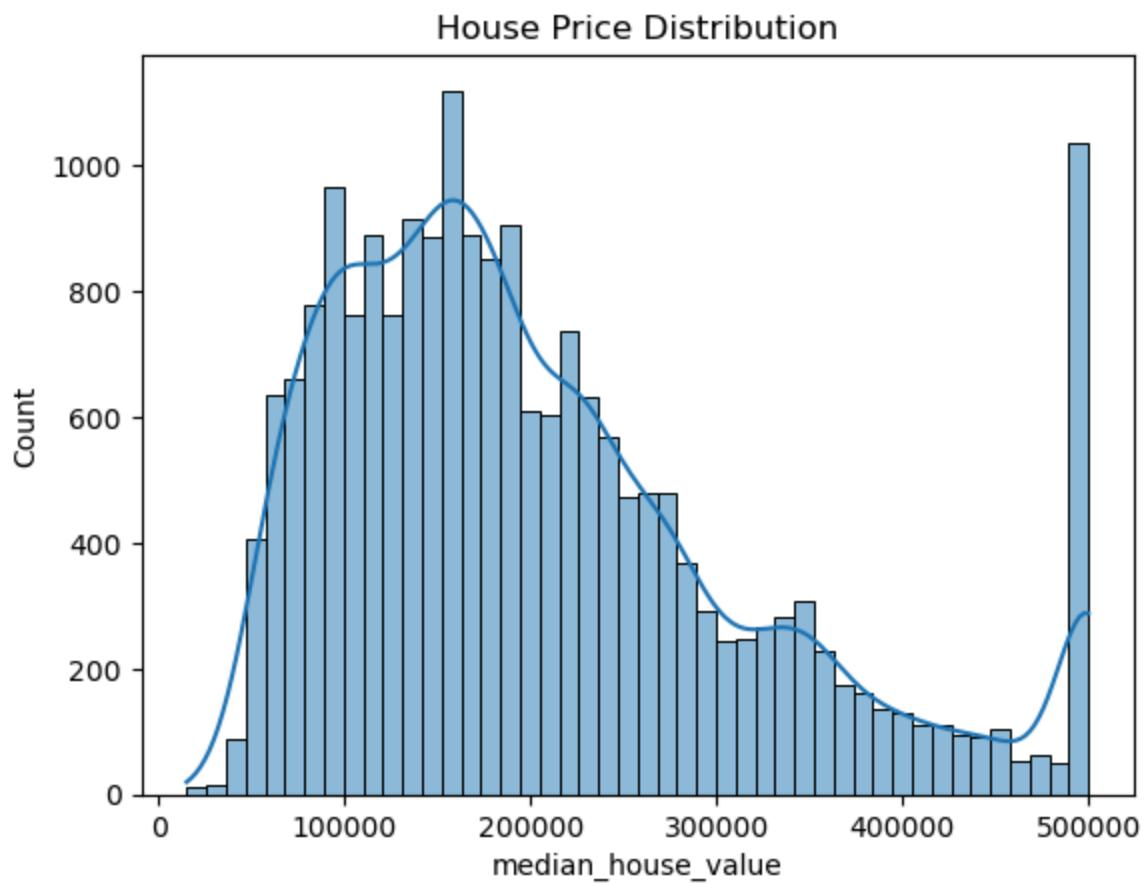


```
In [13]: housing.isnull().sum()
```

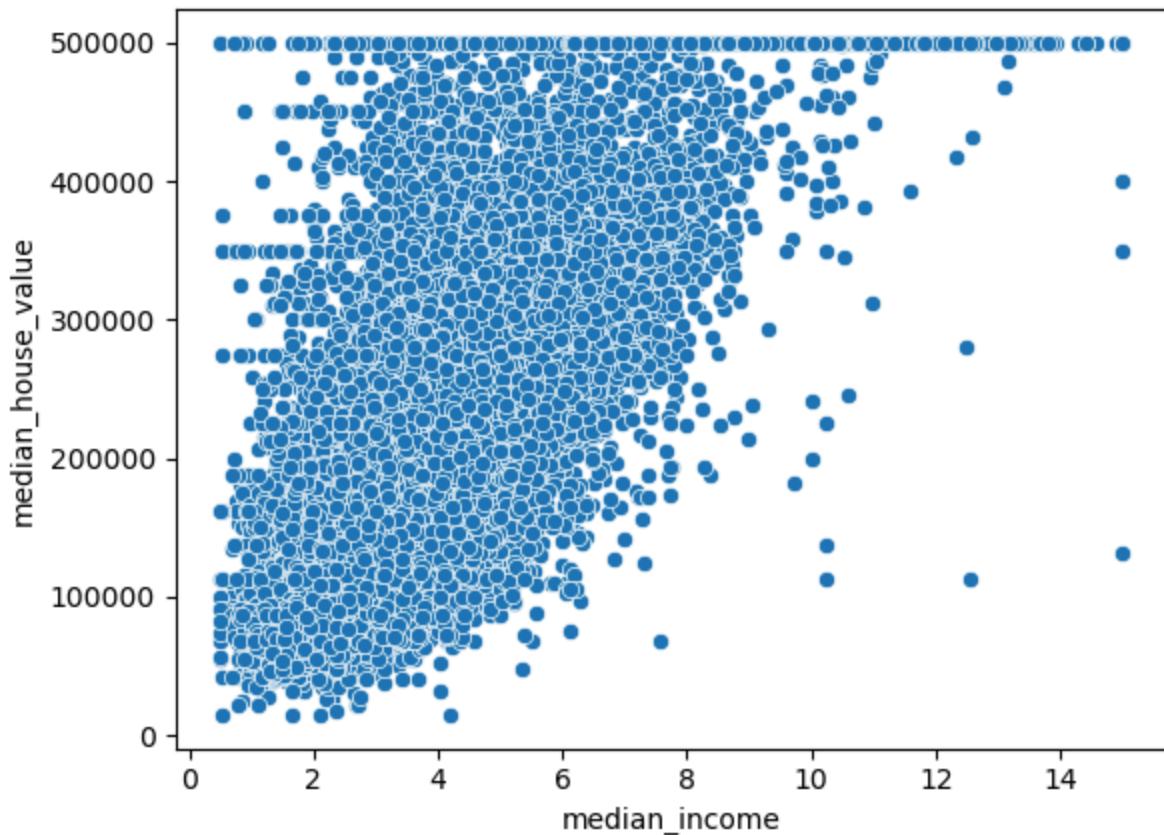
```
Out[13]: longitude      0
          latitude       0
          housing_median_age 0
          total_rooms      0
          total_bedrooms   0
          population       0
          households       0
          median_income    0
          median_house_value 0
          ocean_proximity  0
          dtype: int64
```

1. Distribution of target variable

```
In [14]: sns.histplot(housing["median_house_value"], kde=True)
plt.title("House Price Distribution")
plt.show()
```



```
In [15]: # using Sactter Plot  
sns.scatterplot(data=housing, x="median_income", y="median_house_value")  
plt.show()
```



Feature Engineering

a. Define X and y

```
In [16]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline

from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.ensemble import RandomForestRegressor

from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
In [17]: X = housing.drop('median_house_value', axis = 1)
y = housing['median_house_value']
```

```
In [18]: X
```

Out[18]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	...
0	-122.23	37.88	41.0	880.0	129.0	
1	-122.22	37.86	21.0	7099.0	1106.0	
2	-122.24	37.85	52.0	1467.0	190.0	
3	-122.25	37.85	52.0	1274.0	235.0	
4	-122.25	37.85	52.0	1627.0	280.0	
...
20635	-121.09	39.48	25.0	1665.0	374.0	
20636	-121.21	39.49	18.0	697.0	150.0	
20637	-121.22	39.43	17.0	2254.0	485.0	
20638	-121.32	39.43	18.0	1860.0	409.0	
20639	-121.24	39.37	16.0	2785.0	616.0	

20640 rows × 9 columns

In [19]: y

Out[19]:

```
0      452600.0
1      358500.0
2      352100.0
3      341300.0
4      342200.0
      ...
20635    78100.0
20636    77100.0
20637    92300.0
20638    84700.0
20639    89400.0
Name: median_house_value, Length: 20640, dtype: float64
```

b. Column Separation

In [20]:

```
numeric_cols = X.select_dtypes(include=['int64', 'float64']).columns
cat_cols = ["ocean_proximity"]
```

In [21]: numeric_cols

Out[21]:

```
Index(['longitude', 'latitude', 'housing_median_age', 'total_rooms',
       'total_bedrooms', 'population', 'households', 'median_income'],
      dtype='object')
```

In [22]: cat_cols

Out[22]:

```
['ocean_proximity']
```

c. Preprocessing PipeLine

```
In [23]: preprocessor = ColumnTransformer(  
    transformers=[  
        ("num", StandardScaler(), numeric_cols),  
        ("cat", OneHotEncoder(handle_unknown='ignore'), cat_cols)  
    ]  
)
```

Traing and Testing Test

```
In [24]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, ran
```

```
In [25]: X_train
```

```
Out[25]:      longitude  latitude  housing_median_age  total_rooms  total_bedrooms  ...  
  14196     -117.03    32.71            33.0       3126.0        627.0  
  8267     -118.16    33.77            49.0       3382.0        787.0  
 17445     -120.48    34.66             4.0       1897.0        331.0  
 14265     -117.11    32.69            36.0       1421.0        367.0  
 2271     -119.80    36.78            43.0       2382.0        431.0  
  ...         ...       ...           ...          ...        ...  
 11284     -117.96    33.78            35.0       1330.0        201.0  
 11964     -117.43    34.02            33.0       3084.0        570.0  
 5390     -118.38    34.03            36.0       2101.0        569.0  
 860      -121.96    37.58            15.0       3575.0        597.0  
 15795    -122.42    37.77            52.0       4226.0        1315.0
```

16512 rows × 9 columns

```
In [26]: X_test
```

Out[26]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	...
20046	-119.01	36.06	25.0	1505.0	435.0	
3024	-119.46	35.14	30.0	2943.0	435.0	
15663	-122.44	37.80	52.0	3830.0	435.0	
20484	-118.72	34.28	17.0	3051.0	435.0	
9814	-121.93	36.62	34.0	2351.0	435.0	
...
15362	-117.22	33.36	16.0	3165.0	482.0	
16623	-120.83	35.36	28.0	4323.0	886.0	
18086	-122.05	37.31	25.0	4111.0	538.0	
2144	-119.76	36.77	36.0	2507.0	466.0	
3665	-118.37	34.22	17.0	1787.0	463.0	

4128 rows × 9 columns

In [27]: y_train

Out[27]:

```
14196    103000.0
8267     382100.0
17445    172600.0
14265    93400.0
2271     96500.0
...
11284    229200.0
11964    97800.0
5390     222100.0
860      283500.0
15795    325000.0
Name: median_house_value, Length: 16512, dtype: float64
```

In [28]: y_test

Out[28]:

```
20046    47700.0
3024     45800.0
15663    500001.0
20484    218600.0
9814     278000.0
...
15362    263300.0
16623    266800.0
18086    500001.0
2144     72300.0
3665     151500.0
Name: median_house_value, Length: 4128, dtype: float64
```

Models Training

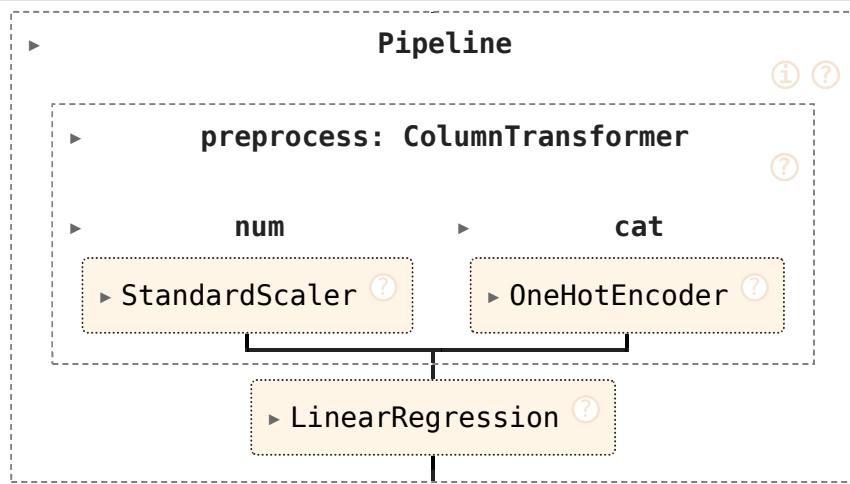
```
In [ ]: model = {}
```

1. Linear regression

```
In [29]: lr_model = Pipeline(steps = [
    ('preprocess', preprocessor),
    ('model', LinearRegression())
])
```

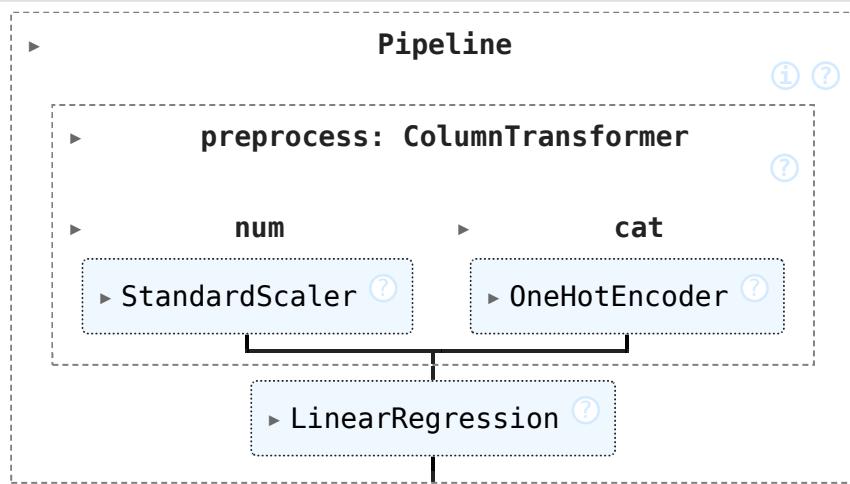
```
In [30]: lr_model
```

```
Out[30]:
```



```
In [45]: lr_model.fit(X_train, y_train)
```

```
Out[45]:
```



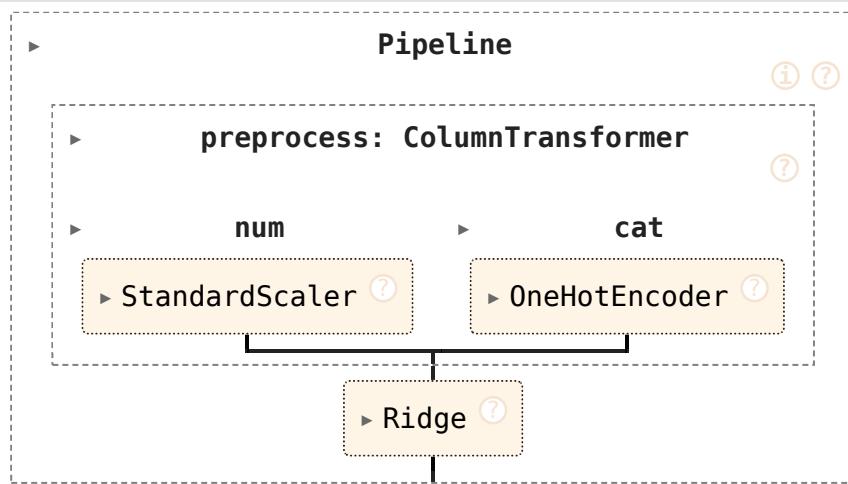
```
In [46]: models["Linear Regression"] = lr_model
```

2. Ridge Regression

```
In [31]: ridge_model = Pipeline(steps=[  
    ('preprocess', preprocess),  
    ('model', Ridge(alpha = 1.0))  
])
```

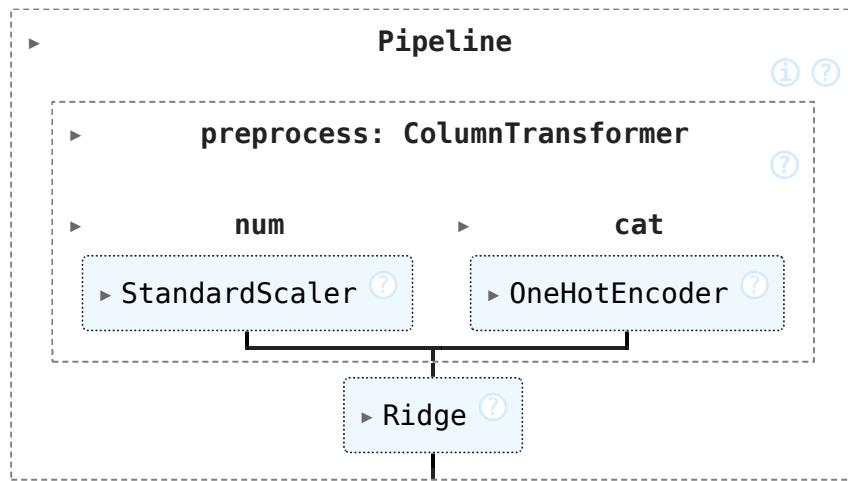
```
In [32]: ridge_model
```

```
Out[32]:
```



```
In [47]: ridge_model.fit(X_train, y_train)
```

```
Out[47]:
```



```
In [48]: models["Ridge Regression"] = ridge_model
```

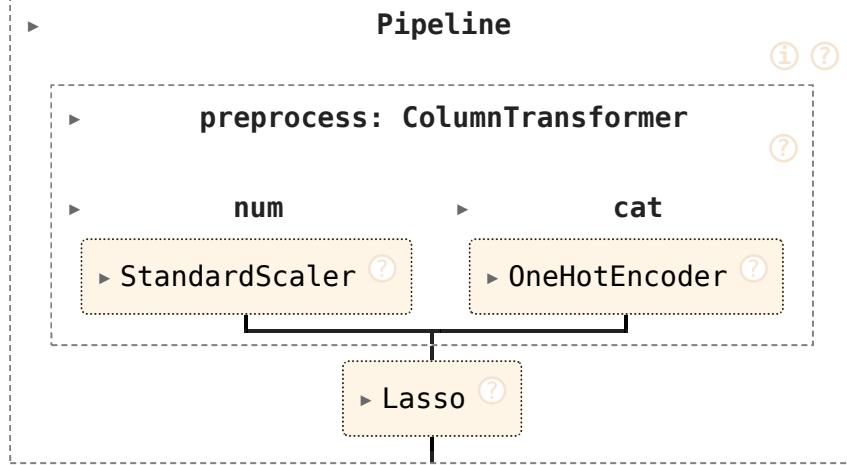
3. Lasso Regression

```
In [33]: lasso_model = Pipeline(steps=[  
    ('preprocess', preprocess),  
    ('model', Lasso(alpha=0.0001))
```

```
])
```

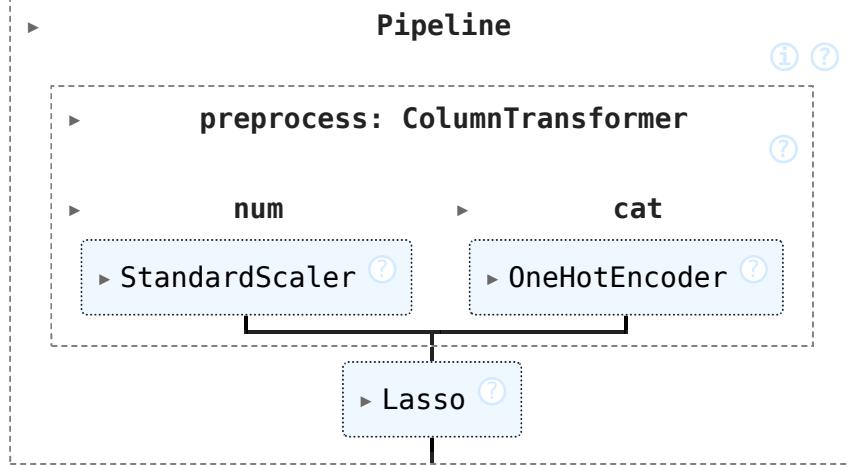
```
In [34]: lasso_model
```

```
Out[34]:
```



```
In [35]: lasso_model.fit(X_train, y_train)
```

```
Out[35]:
```



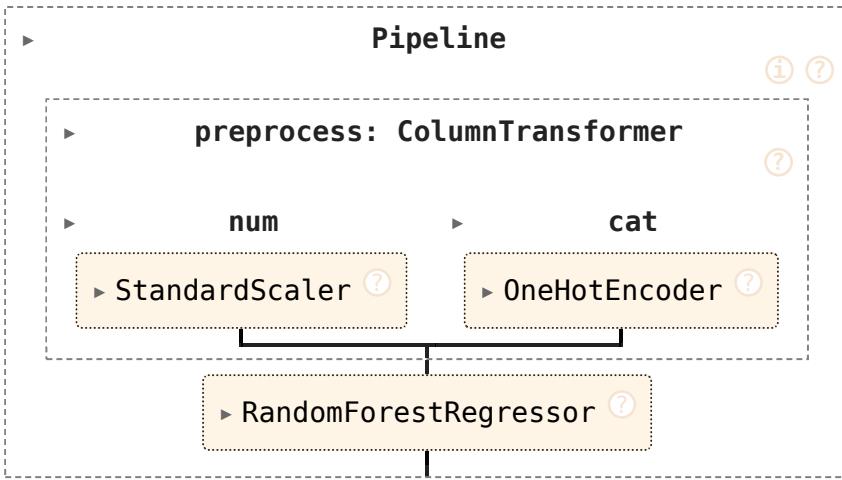
```
In [49]: models["Lasso Regression"] = lasso_model
```

4. Random Forest Regressor

```
In [36]: rf_model = Pipeline(steps=[  
    ('preprocess', preprocess),  
    ('model', RandomForestRegressor(n_estimators=200, random_state=42))  
])
```

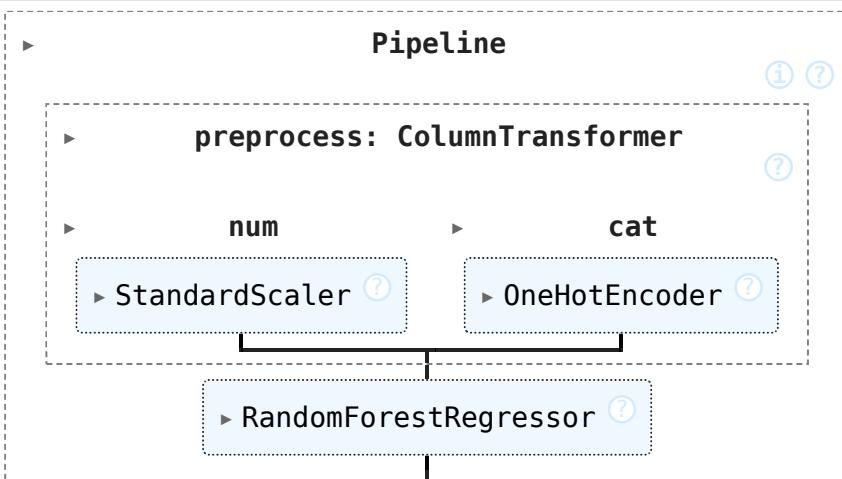
```
In [37]: rf_model
```

```
Out[37]:
```



```
In [38]: rf_model.fit(X_train, y_train)
```

```
Out[38]:
```



```
In [50]: models["Random Forest"] = rf_model
```

MODEL EVALUATION FUNCTION

```
In [39]: def evaluate_model(model, X_test, y_test):
    y_pred = model.predict(X_test)
    mae = mean_absolute_error(y_test, y_pred)
    mse = mean_squared_error(y_test, y_pred)
    rmse = np.sqrt(mse)
    r2 = r2_score(y_test, y_pred)

    return mae, mse, rmse, r2
```

CHECKING MODEL SCORES

```
In [51]: # models = {
```

```

# "Linear Regression": lr_model,
# "Ridge Regression": ridge_model,
# "Lasso Regression": lasso_model,
# "Random Forest": rf_model
#}

results = {}

for name, model in models.items():
    mae, mse, rmse, r2 = evaluate_model(model, X_test, y_test)
    results[name] = [mae, mse, rmse, r2]

results_df = pd.DataFrame(results, index=["MAE", "MSE", "RMSE", "R2 Score"])
results_df

```

Out[51]:

	Linear Regression	Ridge Regression	Lasso Regression	Random Forest
MAE	5.067074e+04	5.067717e+04	5.067074e+04	3.146493e+04
MSE	4.908477e+09	4.909433e+09	4.908477e+09	2.379640e+09
RMSE	7.006052e+04	7.006735e+04	7.006052e+04	4.878156e+04
R2 Score	6.254241e-01	6.253511e-01	6.254241e-01	8.184047e-01

Conclusion

In [52]:

```
print("Best Model Based on R2 Score:")
print(results_df.loc["R2 Score"].idxmax())
```

Best Model Based on R2 Score:
Random Forest

Model Comparison Chart

In [55]:

```
# Extract R2 scores from results table
r2_scores = results_df.loc["R2 Score"]

plt.figure(figsize=(10,6))
plt.bar(r2_scores.index, r2_scores.values)
plt.ylabel("R2 Score")
plt.title("Model Comparison (R2 Score)")
plt.xticks(rotation=45)
plt.show()
```

Model Comparison (R2 Score)

